

Lab 2: Parallel Computing

Introduction to Statistical Computing

Elena Conderana & Sergio Cuenca

Contents

Installing and loading packages	2
Q1. Is parallelization worth it?	2

This lab is to be done outside of class time. You may collaborate with one classmate, but you must identify yourself and his/her name above, in the author's field, and you must submit **your own** lab as this completed .Rmd file.

Installing and loading packages

In order to perform the exercise in this practice you should install and load the `doParallel` package.

```
library(doParallel)
```

```
## Warning: package 'doParallel' was built under R version 4.3.3
```

```
## Loading required package: foreach
```

```
## Warning: package 'foreach' was built under R version 4.3.3
```

```
## Loading required package: iterators
```

```
## Warning: package 'iterators' was built under R version 4.3.3
```

```
## Loading required package: parallel
```

Q1. Is parallelization worth it?

1a. First, let's check how many cores do you have in your computer. Create a new variable `no_cores` equal to the number of cores minus 1.

```
no_cores <- parallel::detectCores()-1
```

1b. Register the cores and prepare the clusters using the `registerDoParallel()` and `makeCluster()` function. This will allow to parallelize code in the following code chunks.

```
registerDoParallel(no_cores)
cl <- makeCluster(no_cores)
```

1c. Now, you have the following function which calculates the prime numbers from 1 to `n`. Use the `microbenchmark` package to check which is faster to calculate the prime numbers when `n` goes from 10 to 10000: `lapply`, a `for` loop, `parLapply` or a `foreach` loop. The `lapply` and the `for` loop have been written for you. Which function is faster?

```
library(microbenchmark)
```

```
## Warning: package 'microbenchmark' was built under R version 4.3.3
```

```

getPrimeNumbers <- function(n) {
  n <- as.integer(n)
  if(n > 1e6) stop("n too large")
  primes <- rep(TRUE, n)
  primes[1] <- FALSE
  last.prime <- 2L
  for(i in last.prime:floor(sqrt(n)))
  {
    primes[seq.int(2L*last.prime, n, last.prime)] <- FALSE
    last.prime <- last.prime + min(which(primes[(last.prime+1):n]))
  }
  which(primes)
}
n_vec <- 10:10000
lapplyPrimeNumbers <- function(n_vec) {
  result <- lapply(n_vec, getPrimeNumbers)
}
forPrimeNumbers <- function(n_vec) {
  result <- list()
  for (n in n_vec) {
    result[[n]] <- getPrimeNumbers(n)
  }
}

parLapplyPrimeNumbers <- function(n_vec, cl) {
  clusterExport(cl, "getPrimeNumbers")
  result <- parLapply(cl, n_vec, getPrimeNumbers)
}

foreachPrimeNumbers <- function(n_vec) {
  result <- foreach(n = n_vec, .export = "getPrimeNumbers") %dopar% {
    getPrimeNumbers(n)
  }
}

bench_results <- microbenchmark(
  for_loop = forPrimeNumbers(n_vec),
  lapply = lapplyPrimeNumbers(n_vec),
  parLapply = parLapplyPrimeNumbers(n_vec, cl),
  foreach = foreachPrimeNumbers(n_vec),
  times = 10
)

print(bench_results)

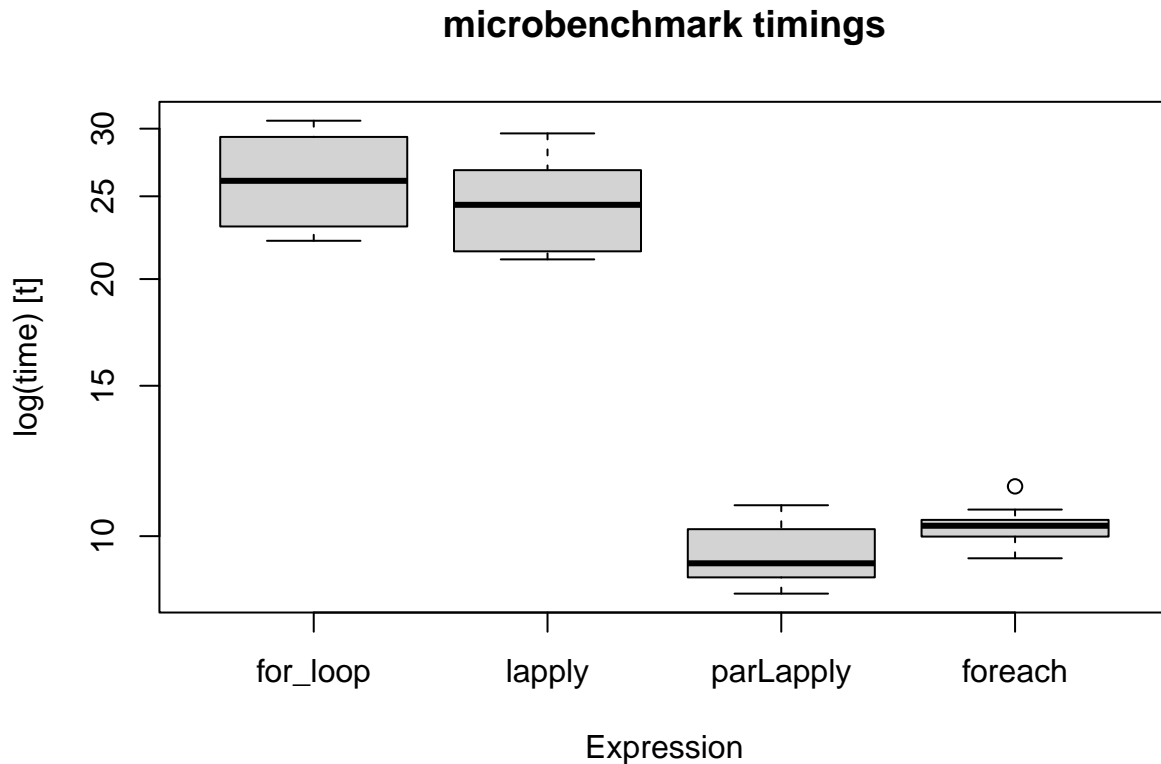
```

```

## Unit: seconds
##      expr      min       lq      mean     median        uq      max neval
## for_loop 22.173259 23.046202 26.24267 26.065240 29.33905 30.64889    10
## lapply   21.090017 21.552563 24.63915 24.431811 26.82444 29.61922    10
## parLapply 8.565109 8.947913 9.53242 9.294351 10.19034 10.86899    10
## foreach  9.420767 9.988525 10.26304 10.284520 10.44830 11.43652    10

```

```
boxplot(bench_results)
```



1d Remember to use stop the clusters in `cl` using the `stopCluster` function.

```
stopCluster(cl)
```

Challenge 01. Search around your computer for a sequential code that might be parallelized. Using the `doParallel` package, parallelize the code and calculate the speedup. If you cannot find any code to parallelize, use the following code:

```
x <- iris[which(iris[,5] != "setosa"), c(1,5)]
trials <- seq(1, 10000)
boot_fx <- function(trial) {
  ind <- sample(100, 100, replace=TRUE)
  result1 <- glm(x[ind,2]~x[ind,1], family=binomial(logit))
  r <- coefficients(result1)
  res <- rbind(data.frame(), r)
}

cl <- makeCluster(no_cores)
clusterExport(cl, c("boot_fx", "x"))

parLapplyBoot <- function(trial, cl) {
  result <- parLapply(cl, trial, boot_fx)
}
```

```

bench_results_boot <- microbenchmark(
  boot_fx = lapply(trials, boot_fx),
  parLapplyBoot = parLapplyBoot(trials, cl),
  times = 10
)

```

```

print(bench_results_boot)

```

```

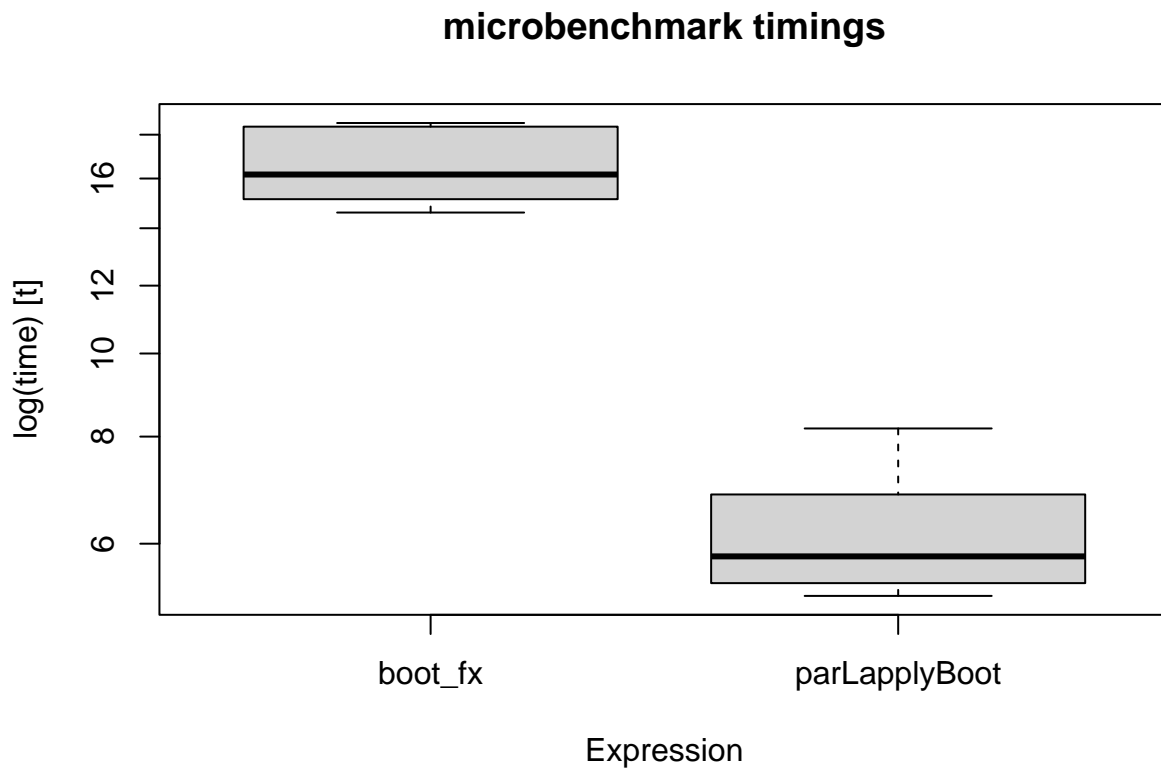
## Unit: seconds
##      expr      min       lq      mean   median       uq      max
##  boot_fx 14.602741 15.135870 16.547065 16.173434 18.392419 18.573014
## parLapplyBoot 5.215137 5.397126 6.181511 5.799208 6.848456 8.176862
##   neval
##      10
##      10

```

```

boxplot(bench_results_boot)

```



```

stopCluster(cl)

```